
Smoothing Binary Optimization: A Primal-Dual Perspective

Wenbo Liu¹, Akang Wang^{2,3}, Dun Ma¹, Hongyi Jiang⁴, Jianghua Wu², and Wenguo Yang¹

¹University of Chinese Academy of Sciences, China

²Shenzhen Research Institute of Big Data, China

³The Chinese University of Hong Kong, Shenzhen, China

⁴City University of Hong Kong, Hong Kong SAR, China

Abstract

Binary optimization is a powerful tool for modeling combinatorial problems, yet scalable and theoretically sound solution methods remain elusive. Existing exact and heuristic solvers often face a trade-off between scalability and solution quality on large-scale instances. In this work, we introduce a novel primal-dual framework that reformulates unconstrained binary optimization as a continuous minimax problem, satisfying a strong max-min property. This reformulation effectively smooths the discrete problem, enabling the application of efficient gradient-based methods. We propose a simultaneous gradient descent-ascent algorithm that is highly parallelizable on GPUs and provably converges to a binary feasible solution in sublinear time. Extensive experiments on large-scale problems with up to 20,000 variables demonstrate that our method identifies high-quality solutions within seconds, significantly outperforming state-of-the-art alternatives.

1 Introduction

In this work, we consider the following *binary optimization* problem:

$$\min_{\mathbf{x} \in \{0,1\}^n} F(\mathbf{x}), \quad (1)$$

where $\mathbf{x}$ denotes a binary variable of dimension n and $F : \{0,1\}^n \rightarrow \mathbb{R}$ represents an arbitrary real-valued function defined on the binary domain. Problem (1) serves as a general model for a wide range of combinatorial optimization problems. A prominent special case is the *quadratic unconstrained binary optimization* (QUBO), which is inherently equivalent to the Max-Cut problem and capable of characterizing all of the Karp's 21 $\mathcal{NP}$ -complete problems [1]. Moreover, *integer linear programs* can be approximately transformed into QUBO by appropriately penalizing the constraints via quadratic terms [2]. Other classic problems such as *maximum satisfiability* (MaxSAT) [3] are also encompassed within this framework.

Exact methods for binary optimization typically rely on branch-and-bound frameworks, integrated with various relaxation techniques such as the linear programming-based relaxations [4] and the semi-definite relaxations [5]. Recent advances have further enhanced these methods through sophisticated problem reduction techniques and accelerated cutting plane separation [6]. However, the inherent $\mathcal{NP}$ -hardness fundamentally limits the scalability of exact approaches, which generally remain practical only for instances involving up to a few hundred variables.

Given the limitations of exact methods, research has shifted toward heuristics for large-scale binary optimization. Traditional CPU-based methods involving customized tabu search [7, 8], simulated annealing [9] and genetic algorithms [10], rely on well-crafted search mechanisms to iteratively

explore the solution space. However, their performance is often constrained by sequential execution. Recent hardware progress has enabled significant improvements through GPU-acceleration, and conventional heuristics have benefited. For instance, ABS2 [11] integrates multiple search strategies and genetic operations executed in parallel on GPUs, while FEM [12] introduces the simulated annealing method from a continuous view and optimized via GPU implementations. GPU acceleration has also propelled the development of first-order methods, yielding frameworks such as the sharp-peak penalized continuous relaxation by [13] and the sampling-based gradient policy by [14]. Beyond classical computing approaches, quantum annealing on specialized hardware [15] has emerged as a promising alternative for QUBO. However, the limited number of qubits and noise sensitivity [16] still impede its widespread practical application.

Meanwhile, learning-based methods are gaining substantial traction as a powerful approach to binary optimization. For example, QUBO problems can be naturally represented using graph structures, enabling the application of *graph neural networks* (GNNs) to effectively model variable interactions and achieve strong empirical performance [17, 18]. In addition, [19] introduces a general graph representation for *constraint satisfaction problems* and proposes a GNN-based reinforcement learning approach enhanced with global search actions. Despite their promise, such methods often operate without theoretical guarantees and encounter practical difficulties related to data acquisition and model generalization.

A prominent line of work within this domain *smooths* Problem (1) using a probabilistic framework [20, 21]: GNNs are leveraged to parameterize an independent distribution over the binary space $\{0, 1\}^n$, and the loss function is subsequently constructed in terms of the expectation of F . We note that this approach is essentially to relax Problem (1) to:

$$\min_{\mathbf{x} \in [0,1]^n} f(\mathbf{x}), \quad (2)$$

where $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is the *multilinear extension* of F (see Section 2 for definition). Although this relaxation preserves theoretical guarantees relative to the original problem [22], the non-convex nature of f in Problem (2) still presents significant optimization challenges.

To tackle this challenge, we revisit the binary optimization Problem (1), and reformulate it into a continuous model by characterizing binarity via continuous variables and convex constraint functions. Based on Lagrangian duality, an equivalent *minimax* problem is derived, to which we develop a *gradient descent-ascent* (GDA)-based algorithm for efficient solutions. We remark that by initializing dual variables as positive values, we induce convexity in the primal space, promoting the identification of high-quality solutions. As optimization proceeds, the dual variables gradually decrease, shifting emphasis from convexity promotion to the enforcement of integrality. We refer to this approach as “a Primal-Dual approach for Binary Optimization” (denoted as PDBO), with the overall framework illustrated in Figure 1.




Figure 1: An illustration of PDBO.

The distinct contributions of our work are summarized as follows.

- **Primal-Dual Reformulation.** We introduce a novel framework that reformulates unconstrained binary optimization as a continuous minimax problem, satisfying a strong max-min property. This reformulation smooths the discrete problem, enabling efficient gradient-based optimization.
- **GDA with Guarantees.** We develop a simultaneous GDA-based algorithm that is highly parallelizable on GPUs and converges to binary solutions in sublinear time.

- **Empirical Superiority.** Through extensive experiments on large-scale problems including Max-Cut, MaxSAT with up to 20,000 variables, we demonstrate that our method achieves state-of-the-art performance, obtaining high-quality solutions within seconds.

2 Preliminaries

This section establishes foundational definitions—chief among them the *multilinear extension*—and reviews its key properties for later use.

Definition 2.1 (Extension). Let $F : \{0, 1\}^n \rightarrow \mathbb{R}$ be a function defined on the binary domain. A function $f : [0, 1]^n \rightarrow \mathbb{R}$ is called an *extension* of F if $f(\mathbf{x}) = F(\mathbf{x})$ for every $\mathbf{x} \in \{0, 1\}^n$.

Definition 2.2 (Multilinear Extension). For any function $F : \{0, 1\}^n \rightarrow \mathbb{R}$, its multilinear extension f is defined by:

$$f(\mathbf{x}) := \mathbb{E}_{\boldsymbol{\xi} \sim P_{\mathbf{x}}} [F(\boldsymbol{\xi})] = \sum_{\boldsymbol{\xi} \in \{0, 1\}^n} P_{\mathbf{x}}(\boldsymbol{\xi}) \cdot F(\boldsymbol{\xi}),$$

where for $\mathbf{x} \in [0, 1]^n$, $P_{\mathbf{x}}(\boldsymbol{\xi}) := \prod_{i=1}^n x_i^{\xi_i} (1 - x_i)^{1-\xi_i}$ denotes the n -dimensional independent Bernoulli distribution with mean vector $\mathbf{x}$.

We remark that f is affine in each variable when all other variables are held fixed. A key property of the multilinear extension is that it preserves the minimum value of the original function:

Lemma 2.3. Let $f : [0, 1]^n \rightarrow \mathbb{R}$ be the multilinear extension of $F : \{0, 1\}^n \rightarrow \mathbb{R}$. Then

$$\min_{\mathbf{x} \in [0, 1]^n} f(\mathbf{x}) = \min_{\mathbf{x} \in \{0, 1\}^n} F(\mathbf{x}).$$

Lemma 2.3 establishes the equivalence between Problem (1) and its continuous relaxation (2). However, we note that f is generally non-convex, making the relaxed problem still challenging to solve.

Proposition 2.4. If $F : \{0, 1\}^n \rightarrow \mathbb{R}$ is a polynomial function of the form $F(\mathbf{x}) = \sum_{\mathbf{d}} a_{\mathbf{d}} \mathbf{x}^{\mathbf{d}}$, where $\mathbf{x}^{\mathbf{d}} := x_1^{d_1} x_2^{d_2} \cdots x_n^{d_n}$, then its multilinear extension is given by:

$$f(\mathbf{x}) = \sum_{\mathbf{d}} a_{\mathbf{d}} \prod_{i: d_i \neq 0} x_i.$$

Proposition 2.4 shows that for polynomial objectives (such as QUBO and MaxSAT), the multilinear extension has a compact closed form, avoiding exponential summation. This makes computation tractable for many practical optimization problems.

Definition 2.5 (Minimax Problem and Critical Points). Consider the minimax problem

$$\min_{\mathbf{x} \in \mathcal{X}} \max_{\mathbf{y} \in \mathcal{Y}} L(\mathbf{x}, \mathbf{y}),$$

where $\mathcal{X} \subseteq \mathbb{R}^n$ and $\mathcal{Y} \subseteq \mathbb{R}^m$ are the domains of $\mathbf{x}$ and $\mathbf{y}$, respectively. For a point $(\mathbf{x}^*, \mathbf{y}^*) \in \mathcal{X} \times \mathcal{Y}$:

- $(\mathbf{x}^*, \mathbf{y}^*)$ is a *saddle point* if for all $(\mathbf{x}, \mathbf{y}) \in \mathcal{X} \times \mathcal{Y}$: $L(\mathbf{x}^*, \mathbf{y}) \leq L(\mathbf{x}^*, \mathbf{y}^*) \leq L(\mathbf{x}, \mathbf{y}^*)$.
- $(\mathbf{x}^*, \mathbf{y}^*)$ is a *stationary point* if $\nabla_{\mathbf{x}} L(\mathbf{x}^*, \mathbf{y}^*) = 0$ and $\nabla_{\mathbf{y}} L(\mathbf{x}^*, \mathbf{y}^*) = 0$.

Remark 2.6. A saddle point does not necessarily exist for an arbitrary minimax problem. However, if the strong max-min property holds, i.e.,

$$\min_{\mathbf{x} \in \mathcal{X}} \max_{\mathbf{y} \in \mathcal{Y}} L(\mathbf{x}, \mathbf{y}) = \max_{\mathbf{y} \in \mathcal{Y}} \min_{\mathbf{x} \in \mathcal{X}} L(\mathbf{x}, \mathbf{y}) \in \mathbb{R},$$

then there exists at least one saddle point.

3 Methodology

In Section 3.1, we introduce a primal-dual framework that reformulates the binary optimization Problem (1) as a minimax problem. Then, in Section 3.2, we develop a GDA-based algorithm to solve the resulting minimax formulation and establish its convergence guarantees.

3.1 Primal-Dual Reformulation

We begin by reformulating Problem (1) as a constrained continuous optimization problem:

$$\begin{aligned} \min_{\mathbf{x} \in [0,1]^n} \quad & f(\mathbf{x}) \\ \text{s.t.} \quad & g(x_i) = 0 \quad \forall i \in [n], \end{aligned} \quad (3)$$

where f is the multilinear extension of F , and $g : [0, 1] \rightarrow \mathbb{R}$ is a constraint function to enforce binarity such that the condition $g(x_i) = 0$ is equivalent to $x_i \in \{0, 1\}$.

An appropriate function g should satisfy the following sufficient conditions:

- g is convex and continuous on $[0, 1]$, with $g(0) = g(1) = 0$;
- g is differentiable on $(0, 1)$, and $g'(x) = 0$ if and only if $x = \frac{1}{2}$;
- g is symmetric on $[0, 1]$, i.e., $g(x) = g(1 - x)$ for all $x \in [0, 1]$.

These conditions imply that $g(x) < 0$ for all $x \in (0, 1)$ and $g'(x) < 0$ for all $x \in (0, \frac{1}{2})$, and therefore $-g$ can be regarded as a metric of fractionality. Several choices of g have been proposed in the literature. Common examples include an even-degree polynomial $g(x) := (2x - 1)^{2d} - 1$ [18] and the entropy-based function $g(x) := x \log(x) + (1 - x) \log(1 - x)$ [12], extended continuously so that $g(0) = g(1) = 0$.

Let $\mathbf{y} \in \mathbb{R}^n$ denote the dual variables associated with the constraints in Problem (3). We define the Lagrangian function as:

$$L(\mathbf{x}, \mathbf{y}) := f(\mathbf{x}) + \sum_{i=1}^n y_i g(x_i).$$

Proposition 3.1. *Let v^* denote the optimal value of Problem (3), then $v^* = \inf_{\mathbf{x} \in [0,1]^n} \sup_{\mathbf{y} \in \mathbb{R}^n} L(\mathbf{x}, \mathbf{y})$.*

Proposition 3.1 characterizes the primal problem as a minimax optimization. The following theorem establishes strong duality for Problem (3).

Theorem 3.2.

$$\inf_{\mathbf{x} \in [0,1]^n} \sup_{\mathbf{y} \in \mathbb{R}^n} L(\mathbf{x}, \mathbf{y}) = \sup_{\mathbf{y} \in \mathbb{R}^n} \inf_{\mathbf{x} \in [0,1]^n} L(\mathbf{x}, \mathbf{y}). \quad (4)$$

Equation (4) in Theorem 3.2, known as the *strong max-min property* [23], implies the existence of a *saddle point* of the Lagrangian function that corresponds to an optimal solution of Problem (3). We therefore reformulate Problem (3) as the following equivalent minimax problem:

$$\min_{\mathbf{x} \in [0,1]^n} \max_{\mathbf{y} \in \mathbb{R}^n} L(\mathbf{x}, \mathbf{y}), \quad (5)$$

where the order of minimization and maximization can be interchanged.

3.2 Minimax Optimization

The Lagrangian $L(\mathbf{x}, \mathbf{y})$ is linear in $\mathbf{y}$ but not necessarily convex in $\mathbf{x}$. Consequently, solving the minimax problem (5) to global optimality remains $\mathcal{NP}$ -hard. In this section, we develop a GDA-based algorithm for solving the minimax problem (5) and establish its convergence guarantees.

3.2.1 Gradient Descent-Ascent

For Problem (5), a classical approach is the GDA method, which alternates between gradient descent on the primal variable $\mathbf{x}$ and gradient ascent on the dual variable $\mathbf{y}$. The strong max-min property enables *simultaneous* updates as follows:

$$\begin{aligned} \mathbf{x}^{t+1} &\leftarrow \Pi_{[0,1]^n} (\mathbf{x}^t - \alpha \cdot \nabla_{\mathbf{x}} L(\mathbf{x}^t, \mathbf{y}^t)), \\ \mathbf{y}^{t+1} &\leftarrow \mathbf{y}^t + \beta \cdot \nabla_{\mathbf{y}} L(\mathbf{x}^t, \mathbf{y}^t), \end{aligned} \quad (6)$$

where $\mathbf{x}^t$ and $\mathbf{y}^t$ denote the primal and dual variables at iteration t , $\Pi_{[0,1]^n}$ denotes the projection onto the hypercube $[0, 1]^n$, and $\alpha, \beta > 0$ are the step sizes.

Notably, the projection ensures $\mathbf{x}^t \in [0, 1]^n$ throughout the optimization, which implies the non-increasing monotonicity of $\{y_i^t\}_t$ since $g(x_i^t) \leq 0$. Moreover, the term $\sum_{i=1}^n (-y_i) \cdot (-g(x_i))$ in $L(\mathbf{x}, \mathbf{y})$ can be interpreted as an integrality penalty, where the penalty coefficients $-y_i$ increase monotonically during optimization. Moreover, GDA introduces a principled update rule for the penalty coefficients, distinguishing PDBO from traditional penalty methods. However, increasing penalties alone does not guarantee integer solutions, as fractional stationary points may persist.

Example (Max-Cut). Let $f(\mathbf{x}) := \mathbf{x}^\top \mathbf{W} \mathbf{x} - \mathbf{1}^\top \mathbf{W} \mathbf{x}$ and $g(x_i) := x_i^2 - x_i$. Then the fractional solution $\mathbf{x}^t := (\frac{1}{2}, \dots, \frac{1}{2}) \in \mathbb{R}^n$ yields $\nabla_{\mathbf{x}} L(\mathbf{x}^t, \mathbf{y}^t) = 0$ for any $\mathbf{y}^t \in \mathbb{R}^n$.

To overcome this limitation, we incorporate targeted perturbations to prevent convergence to fractional points. Specifically, given an initial solution $(\mathbf{x}^0, \mathbf{y}^0) \in [0, 1]^n \times \mathbb{R}_{++}^n$, maximum iterations $T_{\max}$, and tolerance $\delta > 0$, we iteratively perform the updates in (6). When x_i^t is closed to and tends to stagnate near $\frac{1}{2}$, we perturb it to maintain a minimum distance of δ from $\frac{1}{2}$. The complete procedure is described in Algorithm 1.

Algorithm 1 PDBO

Require: Initial solution $(\mathbf{x}^0, \mathbf{y}^0) \in [0, 1]^n \times \mathbb{R}_{++}^n$, Stepsizes $(\alpha, \beta) \in \mathbb{R}_{++}^2$, Tolerance $\delta > 0$, Number of iterations $T_{\max}$

```

1: for  $t = 0, 1, 2, \dots, T_{\max} - 1$  do
2:   for  $i = 1, 2, \dots, n$  do
3:     if  $|x_i^t - \frac{1}{2}| \leq \delta$  and  $|\frac{\partial}{\partial x_i} L(\mathbf{x}^t, \mathbf{y}^t)| \leq 2\delta$  and  $y_i^t \leq 0$  then
4:        $x_i^{t+1} \leftarrow \begin{cases} \frac{1}{2} - \delta & \text{if } x_i^t \leq \frac{1}{2}, \\ \frac{1}{2} + \delta & \text{otherwise.} \end{cases}$ 
5:     else
6:        $x_i^{t+1} \leftarrow \Pi_{[0,1]} \left( x_i^t - \alpha \cdot \frac{\partial}{\partial x_i} L(\mathbf{x}^t, \mathbf{y}^t) \right)$ 
7:     end if
8:      $y_i^{t+1} \leftarrow y_i^t + \beta \cdot g(x_i^t)$ 
9:   end for
10: end for
```

3.2.2 Convergence Analysis

This section analyzes the convergence of the primal variable $\mathbf{x}$ in Algorithm 1 to a binary point, and provides corresponding complexity guarantees. We begin by establishing a lower bound for y_i^t .

Proposition 3.3. Define $\Theta := \max_{\mathbf{x} \in [0,1]^n} \|\nabla_{\mathbf{x}} f(\mathbf{x})\|_1$ and $b := \frac{1+\Theta}{g'(\frac{1}{2}-\delta)} + (2 + \lceil \frac{1}{2\alpha} \rceil) \cdot \beta \cdot g(\frac{1}{2})$. Then for each $i \in [n]$ and any $t \geq 0$, we have $y_i^t \geq b$.

Note that for any $i \in [n]$, the sequence $\{y_i^t\}_{t \geq 0}$ is monotone non-increasing, and hence the lower bound leads to the convergence of $\{y_i^t\}_{t \geq 0}$. Furthermore, $\{g(x_i^t)\}_{t \geq 0}$ converges to zero since $\beta g(x_i^t) = y_i^{t+1} - y_i^t$.

Corollary 3.4. The sequence $\{\mathbf{y}^t\}_{t \geq 0}$ converges to some $\mathbf{y}^* \in \mathbb{R}^n$ with $y_i^* \geq b$ for any $i \in [n]$.

Corollary 3.5. For any $i \in [n]$, the sequence $\{g(x_i^t)\}_{t \geq 0}$ converges to 0.

We remark that Corollary 3.5 implies the convergence of $\mathbf{x}$ to a binary point. This holds because $\|\nabla_{\mathbf{x}} L(\mathbf{x}^t, \mathbf{y}^t)\|$ is bounded, and by choosing a sufficiently small step size α , oscillations between binary points can be avoided.

We now establish the iteration complexity for Algorithm 1 to reach a binary solution. Since the primary goal is to solve the binary optimization problem (1), we focus on convergence toward integrality rather than traditional stationarity criteria. The concept is formalized as follows:

Definition 3.6. Given Problem (3), a point $\mathbf{x} \in [0, 1]^n$ is called an ϵ -binary point if $-\sum_{i=1}^n g(x_i) \leq \epsilon$.

Theorem 3.7. The iteration complexity for Algorithm 1 to return an ϵ -binary point is bounded by

$$\mathcal{O} \left(\frac{\|\mathbf{y}^0 - \mathbf{y}^*\|_1}{\beta \epsilon} \right).$$

Theorem 3.7 shows that Algorithm 1 converges to a binary feasible point at a sublinear rate. While the convergence analysis establishes binarity rather than objective optimality, Section 4.4 complements it with a landscape analysis that helps explain the superior empirical performance of PDBO.

4 Numerical Experiments

We conduct numerical experiments to evaluate the performance of our proposed PDBO method. We evaluate PDBO on three application domains that fit our problem framework, as summarized in Section 4.1. Our code is available at: <https://anonymous.4open.science/r/PD-Bo/>.

4.1 Applications

We introduce three classic combinatorial optimization problems that align with Problem (1): Max-Cut, Max- k -SAT, and Max- k -Cut. Additional experiments on Maximum Independent Set (MIS) and Maximum a Posteriori (MAP) inference are presented in Appendix D.

Max-Cut. Given a graph $G := (V, E)$ with adjacency matrix $\mathbf{W}$, the Max-Cut problem seeks to partition the vertex set V into two disjoint subsets that maximize the total weight of edges between them. This can be formulated as the following binary optimization problem:

$$\max_{\mathbf{x} \in \{0,1\}^n} \sum_{(i,j) \in E} W_{ij}(x_i + x_j - 2x_i x_j),$$

where $x_i \in \{0,1\}$ indicates the partition assignment of vertex i , and W_{ij} denotes the weight of edge $(i,j) \in E$. The expression $x_i + x_j - 2x_i x_j$ equals 1 if $x_i \neq x_j$ and 0 otherwise.

Max- k -SAT. Given a conjunctive normal form (CNF) formula with m clauses, where each clause contains at most k literals, the goal is to find a truth assignment that maximizes the number of satisfied clauses. This problem can be formulated as the following binary optimization problem:

$$\min_{\mathbf{x} \in \{0,1\}^n} \sum_{j=1}^m \prod_{l_i \in C_j} p(x_i)$$

where $x_i \in \{0,1\}$ denotes the truth assignment of the i -th variable, C_j denotes the j -th clause, and $p(x_i) = x_i$ if literal l_i appears negated in clause C_j and $1 - x_i$ otherwise. The product term $\prod_{l_i \in C_j} p(x_i)$ equals 1 if clause C_j is unsatisfied and 0 if it is satisfied.

Max- k -Cut. The Max- k -Cut problem generalizes Max-Cut by partitioning the vertex set V into k disjoint subsets, maximizing the total weight of edges crossing between different subsets. The discrete formulation is:

$$\begin{aligned} \max_{\mathbf{X} \in \mathbb{R}^{k \times n}} \quad & \sum_{(i,j) \in E} W_{ij} (1 - \mathbf{X}_{:,i}^\top \mathbf{X}_{:,j}) \\ \text{s.t.} \quad & \mathbf{X}_{:,i} \in \{\mathbf{e}^{(1)}, \dots, \mathbf{e}^{(k)}\}, \quad \forall i \in [n] \end{aligned}$$

where $\mathbf{X}_{:,i}$ is a k -dimensional one-hot vector indicating the partition assignment of vertex i since $\mathbf{e}^{(j)}$ denotes the j -th standard basis vector. The expression $1 - \mathbf{X}_{:,i}^\top \mathbf{X}_{:,j}$ equals 1 if vertices i and j belong to different subsets and 0 otherwise. We remark that Max- k -Cut can be incorporated into our primal-dual framework with a natural generalization. Please see Appendix B for technical details.

4.2 Setup

Datasets. For Max-Cut and Max- k -Cut, we use the Gset benchmark [24], a widely recognized dataset comprising graphs with 800 to 20,000 nodes, totaling 71 instances. For Max- k -SAT, we use the datasets from [19], which include 3CNF, 4CNF, and 5CNF formulas, each containing 50 instances. This enables consistent and comparable evaluation with existing work.

Baselines. We consider two types of baseline algorithms: For traditional heuristics with GPU accelerations, we adopt (i) ABS2 [11]: an evolutionary method featuring three diverse strategies: multiple search, multiple genetic operations and multiple solution pools; and (ii) FEM [12]: an annealing-based method with additional techniques tailored for Max-Cut. For the learning-based approaches, we consider the following state-of-the-art methods: (iii) PIGNN [17]: an unsupervised method for QUBO

problems, which leverage physics-inspired GNNs and delivers commendable performance; (iv) ANYCSP [19]: a GNN-based reinforcement learning approach for constraint satisfaction problems, utilizing a compact graph representation and global search actions; (v) CRA [18]: an annealing-based method that control the convexity of objectives, with GNNs leveraged to further enhance the performance; and (vi) ROS [21]: a GNN-guided relax-optimize-and-sample method with fine tuning techniques, designed for Max- k -Cut problems. Finally, we consider the advanced commercial solver (vii) Gurobi 12.0.2 [25].

Model Settings. To enforce binarity in Problem (3), we employ the function $g(x) := x^2 - x$ for PDB0 unless otherwise specified. To fully utilize GPU parallelization, we sample B independent initial solutions and execute PDB0 in parallel from each starting point, retaining the best solution across all runs. Detailed parameter configurations, including the values of B , y^0 , α , and β , are provided in Appendix C, together with a sensitivity analysis of the key hyperparameters.

Evaluation Configuration. All experiments are conducted with a 180-second time limit using identical hardware configuration: a 12th Gen Intel Core i9-12900K CPU and an NVIDIA GeForce RTX 3090 GPU. **For each method, we report the solution time as the earliest time at which it attains its best solution within the 180-second budget.** PDB0 is implemented in JAX 0.6.1 [26].

4.3 Computational Results

4.3.1 Results for Max-Cut

We evaluate PDB0 by measuring both solution quality (number of cut edges) and computational efficiency (runtime in seconds). Our main results focus on the five largest Gset graphs (G67, G70, G72, G77, and G81), each containing over 10,000 nodes, which present significant computational challenges for state-of-the-art methods. Results for smaller graphs are provided in Appendix D.

Table 1: Results on Gset instances for Max-Cut

Method	G67 (n=10k)		G70 (n=10k)		G72 (n=10k)		G77 (n=14k)		G81 (n=20k)	
	Obj $\uparrow$	Time $\downarrow$	Obj $\uparrow$	Time $\downarrow$	Obj $\uparrow$	Time $\downarrow$	Obj $\uparrow$	Time $\downarrow$	Obj $\uparrow$	Time $\downarrow$
PIGNN	5538	23.7	8534	25.2	5588	44.5	7896	42.1	11078	157.6
CRA	5948	53.7	9240	51.7	6058	53.9	8720	75.9	12450	120.4
ROS	6144	1.5	8872	1.8	6148	1.2	8746	2.2	12320	5.2
ANYCSP	6772	39.9	9379	35.7	6816	36.1	9686	53.5	13670	73.5
FEM	6782	2.4	5120	0.2	6824	2.6	9688	4.0	13684	7.5
ABS2	6880	156.3	9510	175.5	6932	172.2	9824	171.1	13850	177.1
Gurobi	6944	51.9	9514	133.3	6990	171.6	9882	175.1	13848	179.8
PDB0	6872	2.5	9537	1.9	6906	2.0	9812	2.1	13852	3.3

Table 1 summarizes the results for all baselines. The ‘‘Obj.’’ column shows the best objective value found by each method within the 180-second limit. The ‘‘Time’’ entry corresponds to the earliest time at which the best solution was found. Note that $\uparrow$ denotes that larger values are better, and $\downarrow$ indicates that smaller values are preferred. The results show that PDB0 exhibits substantially greater efficiency than other baselines, and delivers the best solutions for G70 and G81. Meanwhile, Gurobi finds the best solutions to the other three instances. It is worth noting, as highlighted in [6], that recent versions of Gurobi have seen major improvements on QUBO and Max-Cut problems, likely due to the integration of specialized heuristics.

Figure 2 compares the solution quality throughout the optimization process, plotting the objective value against runtime for two representative instances. Clearly, PDB0 (purple) achieves and maintains the highest objective values shortly after initialization, demonstrating both rapid convergence and superior solution quality for most of the time horizon. Although Gurobi (black) eventually attains a marginally better solution given the full time budget, PDB0 maintains a substantial advantage throughout the majority of the optimization timeline.




Figure 2: The objective value versus runtime for G81

4.3.2 Results for Max- k -SAT.

We evaluate PDBO against baseline algorithms FEM, ANYCSP and Gurobi on Max- k -SAT using the 3CNF, 4CNF, and 5CNF datasets, each containing 50 problem instances. Since the objective in Max- k -SAT is a polynomial of higher degree, we accordingly adopt

$g(x) := x \log x + (1 - x) \log(1 - x)$ in PDBO to enforce binarity, which yields better performance. Performance is measured by the average number of unsatisfied clauses across all instances in each dataset, as summarized in Table 2, where entries marked with “—” indicate memory issues. We find that both FEM and Gurobi struggle on these problems, yielding inferior solutions compared to ANYCSP and PDBO. Moreover, the results show that PDBO achieves competitive or superior performance compared to ANYCSP, while avoiding the substantial training overhead and significantly longer inference time required by the latter.

Table 2: Results on CNF instances for Max- k -SAT

Method	3CNF		4CNF		5CNF	
	Obj ↓	Time ↓	Obj ↓	Time ↓	Obj ↓	Time ↓
FEM	1885.2 $\pm$ 23.9	0.8 $\pm$ 0.1	—	—	—	—
ANYCSP	1583.3 $\pm$ 17.5	123.8 $\pm$ 40.7	1210.9 $\pm$ 12.6	141.6 $\pm$ 25.9	1213.7 $\pm$ 11.4	141.0 $\pm$ 31.9
Gurobi	9322.6 $\pm$ 64.2	0.0 $\pm$ 0.1	9329.2 $\pm$ 69.7	0.0 $\pm$ 0.0	9310.7 $\pm$ 75.4	0.1 $\pm$ 0.0
PDBO	1582.7 $\pm$ 12.0	0.9 $\pm$ 0.1	1147.4 $\pm$ 0.8	1.2 $\pm$ 0.1	976.0 $\pm$ 9.3	2.2 $\pm$ 0.3

4.3.3 Results for Max- k -cut

We evaluate the performance of PDBO against baseline algorithms ROS, ANYCSP, FEM and Gurobi on Gset instances for the Max-3-Cut problem. Following our evaluation protocol, we present results for instances with more than 10,000 nodes in Table 3, with remaining results provided in Appendix D. Table 3 demonstrates that PDBO consistently outperforms all baseline methods across all evaluated Max-3-Cut instances, achieving superior solution quality with competitive runtime efficiency.

Table 3: Results on Gset instances for Max-3-Cut

Method	G67 (n=10k)		G70 (n=10k)		G72 (n=10k)		G77 (n=14k)		G81 (n=20k)	
	Obj ↑	Time ↓	Obj ↑	Time ↓	Obj ↑	Time ↓	Obj ↑	Time ↓	Obj ↑	Time ↓
ROS	7364	3.0	9983	1.9	7435	2.9	10559	5.3	14907	9.7
ANYCSP	7797	55.9	9909	16.2	7906	58.9	11158	84.0	15727	115.0
FEM	7748	3.3	9999	1.4	7835	0.7	11102	4.4	15683	6.1
Gurobi	7855	178.1	9999	0.7	8045	174.6	11101	179.9	15146	179.9
PDBO	8015	6.0	9999	3.3	8111	4.7	11467	5.0	16191	7.9

4.4 Landscape Analysis

We next use Max-Cut to illustrate how the primal-dual dynamics shape the optimization landscape. Recall the minimization form of Max-Cut, $f(\mathbf{x}) = \mathbf{x}^\top \mathbf{W} \mathbf{x} - \mathbf{1}^\top \mathbf{W} \mathbf{x}$, where $\mathbf{W}$ is symmetric with zero diagonal entries, so the multilinear extension coincides with the quadratic objective itself. With $g(x_i) = x_i^2 - x_i$, the associated Lagrangian is $L(\mathbf{x}, \mathbf{y}) = \mathbf{x}^\top \mathbf{W} \mathbf{x} - \mathbf{1}^\top \mathbf{W} \mathbf{x} + \sum_{i=1}^n y_i (x_i^2 - x_i)$ and its Hessian with respect to $\mathbf{x}$ is

$$\nabla_{\mathbf{x}}^2 L(\mathbf{x}, \mathbf{y}) = 2(\mathbf{W} + \text{diag}(\mathbf{y})).$$

This simple form reveals a convex-to-nonconvex landscape evolution.

Convex initialization: PDBO initializes the dual variables as $\mathbf{y}^0 = \bar{y}\mathbf{1}$; under the condition $\bar{y} \geq -\lambda_{\min}(\mathbf{W})$, one has $\mathbf{W} + \text{diag}(\mathbf{y}^0) \succeq 0$, which renders $L(\cdot, \mathbf{y}^0)$ globally convex. The method therefore starts from a single-basin landscape without spurious local minima.

Gradual nonconvexification: The dual update obeys $y_i^{t+1} - y_i^t = \beta((x_i^t)^2 - x_i^t) \leq 0$, so each coordinate of $\mathbf{y}^t$ decreases monotonically. As a result, the diagonal curvature in $\text{diag}(\mathbf{y}^t)$ is progressively weakened, and the landscape shifts from convex to nonconvex as PDBO proceeds.

Binary recovery: The dual term can be rewritten as $y_i(x_i^2 - x_i) = -y_i(x_i - x_i^2)$, where $x_i - x_i^2 \geq 0$ measures the fractionality of x_i . The decrease of y_i increases the coefficient $-y_i$, thereby inducing a stronger penalty against fractional values and driving the iterates toward binary solutions, in agreement with the convergence analysis provided in Section 3.2.2.

Empirical evidence. Table 4 further supports this interpretation through a sensitivity analysis on the initialization magnitude $\bar{y}$. Each entry reports the obtained Max-Cut objective value. Here, “ $\checkmark$ ” indicates that $\mathbf{W} + \bar{y} \cdot \mathbf{I} \succeq 0$, while “ $\times$ ” indicates otherwise. The results show that the objective value improves substantially once $\bar{y}$ is large enough to make the initial landscape convex.

Table 4: Effect of the initial dual magnitude $\bar{y}$ on Max-Cut objective values.

Instance	$\bar{y} = 0(\times)$	$\bar{y} = 1(\times)$	$\bar{y} = 2(\times)$	$\bar{y} = 4(\checkmark)$	$\bar{y} = 6(\checkmark)$	$\bar{y} = 8(\checkmark)$	$\bar{y} = 10(\checkmark)$
G67	5632	6394	6674	6864	6872	6868	6868
G70	8756	9145	9398	9535	9537	9536	9536
G72	5714	6474	6720	6898	6906	6900	6900
G77	8124	9108	9522	9810	9812	9804	9804
G81	11338	12886	13836	13840	13852	13846	13846

4.5 Discussion

Both baseline algorithms, CRA and FEM, implement a form of *continuation methods* [27]. They minimize the penalized objective $f(\mathbf{x}) + \mu \sum_i g(x_i)$ while gradually annealing the penalty parameter μ . Although this resembles a *primal-dual* treatment of the *single-constraint* formulation

$$\min_{\mathbf{x} \in [0,1]^n} f(\mathbf{x}) \quad \text{s.t.} \quad \sum_{i=1}^n g(x_i) = 0,$$

PDBO differs in two fundamental aspects: (i) **principled update rule**: Continuation-based methods rely on pre-defined annealing schedule for μ , whereas PDBO’s dual update follows formally from GDA dynamics; (ii) **richer choices of extensions**: PDBO considers the extensions in $S := \{f_{\mathbf{y}}(\cdot) : \mathbf{y} \in \mathbb{R}^n\}$, while the continuation methods are restricted to $S' := \{f_{\mu}(\cdot) : \mu \in \mathbb{R}\}$ with $f_{\mu}(\mathbf{x}) := f(\mathbf{x}) + \mu \sum_{i=1}^n g(x_i)$. In the QUBO case, the tightest convex extension in the space S corresponds to the SDP relaxation [28], yielding a tighter bound than the eigenvalue relaxation [29] resulted from S' .

Ablation study. Table 5 validates the advantage of the primal-dual perspective. First, even the single-constraint variant (denoted as “PDBO-s”) already outperforms continuation-based baselines (CRA, FEM) across all instances. Second, by lifting to n separate constraints—thereby enriching the dual space—the full PDBO framework yields

Table 5: Ablation study on Max-Cut

Method	G72 (n=10k)		G77 (n=14k)		G81 (n=20k)	
	Obj $\uparrow$	Time $\downarrow$	Obj $\uparrow$	Time $\downarrow$	Obj $\uparrow$	Time $\downarrow$
CRA	6058	53.9	8720	75.9	12450	120.4
FEM	6824	2.6	9688	4.0	13684	7.5
PDBO-s	6880	1.4	9782	1.8	13786	2.5
PDBO	6906	2.0	9812	2.1	13852	3.3

further improvements in solution quality. This progressive gain confirms that both the principled dual update and the expanded constraint formulation contribute to higher solution quality.

5 Conclusions

In this work, we introduce PDBO, a primal-dual framework for unconstrained binary optimization that reformulates the problem as a minimax optimization and solves it with a tailored GDA algorithm. We establish convergence guarantees for PDBO and demonstrate its effectiveness through extensive experiments on public benchmarks. Our framework opens several promising directions for future research. First, while GDA provides strong empirical performance, exploring more sophisticated minimax solvers, e.g., extragradient methods, could yield better solution quality. Second, the primal-dual perspective could be extended to more complex discrete problems with additional constraints, potentially expanding its applicability to broader classes of combinatorial optimization. The demonstrated effectiveness of PDBO suggests that the reformulation of binary optimization as a minimax problem is a fruitful approach, warranting further investigation into both algorithmic improvements and theoretical understanding.

References

- [1] Andrew Lucas. Ising formulations of many np problems. *Frontiers in physics*, 2:5, 2014.
- [2] Edoardo Alessandrini, Sergi Ramos-Calderer, Ingo Roth, Emiliano Traversi, and Leandro Aolita. Alleviating the quantum big-m problem. *npj Quantum Information*, 11(1):125, 2025.
- [3] Fahiem Bacchus, Matti Järvisalo, and Ruben Martins. Maximum satisfiability. In *Handbook of satisfiability*, pages 929–991. IOS Press, 2021.
- [4] Francisco Barahona, Michael Jünger, and Gerhard Reinelt. Experiments in quadratic 0–1 programming. *Mathematical programming*, 44(1):127–137, 1989.
- [5] Svatopluk Poljak and Franz Rendl. Solving the max-cut problem using eigenvalues. *Discrete Applied Mathematics*, 62(1-3):249–278, 1995.
- [6] Daniel Rehfeldt, Thorsten Koch, and Yuji Shinano. Faster exact solution of sparse maxcut and qubo problems. *Mathematical Programming Computation*, 15(3):445–470, 2023.
- [7] Gintaras Palubeckis. Iterated tabu search for the unconstrained binary quadratic optimization problem. *Informatica*, 17(2):279–296, 2006.
- [8] Yang Wang, Zhipeng Lü, Fred Glover, and Jin-Kao Hao. Probabilistic grasp-tabu search algorithms for the ubqp problem. *Computers & Operations Research*, 40(12):3100–3107, 2013.
- [9] Dimitris Bertsimas and John Tsitsiklis. Simulated annealing. *Statistical science*, 8(1):10–15, 1993.
- [10] Peter Merz and Bernd Freisleben. Genetic algorithms for binary quadratic programming. In *Proceedings of the genetic and evolutionary computation conference*, volume 1, pages 417–424. Morgan Kaufmann Orlando, FL, 1999.
- [11] Koji Nakano, Daisuke Takafuji, Yasuaki Ito, Takashi Yazane, Junko Yano, Shiro Ozaki, Ryota Katsuki, and Rie Mori. Diverse adaptive bulk search: a framework for solving qubo problems on multiple gpus. In *2023 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW)*, pages 314–325. IEEE, 2023.
- [12] Zi-Song Shen, Feng Pan, Yao Wang, Yi-Ding Men, Wen-Biao Xu, Man-Hong Yung, and Pan Zhang. Free-energy machine for combinatorial optimization. *Nature Computational Science*, pages 1–11, 2025.
- [13] Shenglong Zhou, Shuai Li, Hui Zhang, and Ziyang Luo. Sharp-peak functions for exactly penalizing binary integer programming. *arXiv preprint arXiv:2509.00895*, 2025.
- [14] Cheng Chen, Ruitao Chen, Tianyou Li, Ruicheng Ao, and Zaiwen Wen. A monte carlo policy gradient method with local search for binary optimization. *Mathematical Programming*, pages 1–57, 2025.
- [15] Jesse J Berwald. The mathematics of quantum-enabled applications on the d-wave quantum computer. *Not. Am. Math. Soc.*, 66(832):55, 2019.
- [16] Tristan Zaborniak and Rogério de Sousa. Benchmarking hamiltonian noise in the d-wave quantum annealer. *IEEE Transactions on Quantum Engineering*, 2:1–6, 2021.
- [17] Martin JA Schuetz, J Kyle Brubaker, and Helmut G Katzgraber. Combinatorial optimization with physics-inspired graph neural networks. *Nature Machine Intelligence*, 4(4):367–377, 2022.
- [18] Yuma Ichikawa. Controlling continuous relaxation for combinatorial optimization. *Advances in Neural Information Processing Systems*, 37:47189–47216, 2024.
- [19] Jan Tönshoff, Berke Kisin, Jakob Lindner, and Martin Grohe. One model, any csp: Graph neural networks as fast global search heuristics for constraint satisfaction. In *Proceedings of the Thirty-Second International Joint Conference on Artificial Intelligence, IJCAI-23*, pages 4280–4288. International Joint Conferences on Artificial Intelligence Organization, 8 2023. Main Track.

- [20] Nikolaos Karalias and Andreas Loukas. Erdos goes neural: an unsupervised learning framework for combinatorial optimization on graphs. *Advances in Neural Information Processing Systems*, 33:6659–6672, 2020.
- [21] Yeqing Qiu, Xue Ye, Akang Wang, Yiheng Wang, Qingjiang Shi, and Zhi-Quan Luo. Ros: A gnn-based relax-optimize-and-sample framework for max- k -cut problems. In *Forty-second International Conference on Machine Learning*, 2025.
- [22] Haoyu Peter Wang, Nan Wu, Hang Yang, Cong Hao, and Pan Li. Unsupervised learning for combinatorial optimization with principled objective relaxation. *Advances in Neural Information Processing Systems*, 35:31444–31458, 2022.
- [23] Stephen P Boyd and Lieven Vandenbergh. *Convex optimization*. Cambridge university press, 2004.
- [24] Yinyu Ye. The gset dataset. <https://web.stanford.edu/~yyye/yyye/Gset/>, 2003.
- [25] Gurobi Optimization, LLC. Gurobi Optimizer Reference Manual, 2024.
- [26] James Bradbury, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary, Dougal Maclaurin, George Necula, Adam Paszke, Jake VanderPlas, Skye Wanderman-Milne, and Qiao Zhang. JAX: composable transformations of Python+NumPy programs, 2025.
- [27] Axel Séguin and Daniel Kressner. Continuation methods for riemannian optimization. *SIAM Journal on Optimization*, 32(2):1069–1093, 2022.
- [28] Claude Lemaréchal and François Oustry. *Semidefinite relaxations and Lagrangian duality with application to combinatorial optimization*. PhD thesis, INRIA, 1999.
- [29] Carlos J Nohra, Arvind U Raghunathan, and Nikolaos Sahinidis. Spectral relaxations and branching strategies for global optimization of mixed-integer quadratic programs. *SIAM Journal on Optimization*, 31(1):142–171, 2021.
- [30] Jean Barbier, Florent Krzakala, Lenka Zdeborová, and Pan Zhang. The hard-core model on random graphs revisited. In *Journal of Physics: Conference Series*, volume 473, page 012021. IOP Publishing, 2013.
- [31] D Khuê Lê-Huu and Nikos Paragios. Continuous relaxation of map inference: A nonconvex perspective. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 5533–5541, 2018.
- [32] Ariel Jaimovich, Gal Elidan, Hanah Margalit, and Nir Friedman. Towards an integrated protein-protein interaction network: A relational markov network approach. *Journal of Computational Biology*, 13(2):145–164, 2006.
- [33] Jörg H Kappes, Bjoern Andres, Fred A Hamprecht, Christoph Schnörr, Sebastian Nowozin, Dhruv Batra, Sungwoong Kim, Bernhard X Kausler, Thorben Kröger, Jan Lellmann, et al. A comparative study of modern inference techniques for structured discrete energy minimization problems. *International Journal of Computer Vision*, 115(2):155–184, 2015.
- [34] Gal Elidan and A Globerson. The probabilistic inference challenge (pic2011). *Retrieved from*, 2011.
- [35] Yaomin Wang, Chaolong Ying, Xiaodong Luo, and Tianshu Yu. Neurolifting: Neural inference on markov random fields at scale. *arXiv preprint arXiv:2411.18954*, 2024.
- [36] Martin J Wainwright, Tommi S Jaakkola, and Alan S Willsky. Map estimation via agreement on trees: message-passing and linear programming. *IEEE transactions on information theory*, 51(11):3697–3717, 2005.
- [37] Simon De Givry. toulbar2, an exact cost function network solver. In *24ème édition du congrès annuel de la Société Française de Recherche Opérationnelle et d’Aide à la Décision ROADEF 2023*, 2023.

A Proofs

A.1 Proof of Lemma 2.3

Proof. By definition, for any $\mathbf{x} \in [0, 1]^n$, $f(\mathbf{x})$ is the convex combination of values in $\{F(\boldsymbol{\xi}) : \boldsymbol{\xi} \in \{0, 1\}^n\}$, with the weights given by $P_{\mathbf{x}}$. Therefore, $f(\mathbf{x}) \geq \min_{\boldsymbol{\xi} \in \{0, 1\}^n} F(\boldsymbol{\xi})$ for any $\mathbf{x} \in [0, 1]^n$. This means f does not introduce extra minimum on $[0, 1]^n \setminus \{0, 1\}^n$, and hence the conclusion follows evidently. $\square$

A.2 Proof of Proposition 2.4

Proof.

$$\begin{aligned}
f(\mathbf{x}) &= \mathbb{E}_{\boldsymbol{\xi} \sim P_{\mathbf{x}}} \left(\sum_{\mathbf{d}} a_{\mathbf{d}} \xi_1^{d_1} \dots \xi_n^{d_n} \right) && \# \text{by definition} \\
&= \sum_{\mathbf{d}} a_{\mathbf{d}} \left(\mathbb{E}_{\boldsymbol{\xi} \sim P_{\mathbf{x}}} \xi_1^{d_1} \dots \xi_n^{d_n} \right) && \# \text{interchange } \mathbb{E} \text{ and } \sum \\
&= \sum_{\mathbf{d}} a_{\mathbf{d}} \left(\mathbb{E}_{\boldsymbol{\xi} \sim P_{\mathbf{x}}} \xi_1^{d_1} \right) \dots \left(\mathbb{E}_{\boldsymbol{\xi} \sim P_{\mathbf{x}}} \xi_n^{d_n} \right) && \# \xi_i \text{ are independent} \\
&= \sum_{\mathbf{d}} a_{\mathbf{d}} \left(\mathbb{E}_{\boldsymbol{\xi} \sim P_{\mathbf{x}}} \xi_1 \right) \dots \left(\mathbb{E}_{\boldsymbol{\xi} \sim P_{\mathbf{x}}} \xi_n \right) && \# \xi_i^{d_i} = \xi_i \text{ (for } d_i \neq 0), \text{ since } \xi_i \in \{0, 1\} \\
&= \sum_{\mathbf{d}} a_{\mathbf{d}} \prod_{i: d_i \neq 0} x_i && \# \mathbb{E}_{\boldsymbol{\xi} \sim P_{\mathbf{x}}} \xi_i = x_i
\end{aligned}$$

$\square$

A.3 Proof of Proposition 3.1

Proof. Observe that

$$\begin{aligned}
\sup_{\mathbf{y} \in \mathbb{R}^n} L(\mathbf{x}, \mathbf{y}) &= f(\mathbf{x}) + \sup_{\mathbf{y} \in \mathbb{R}^n} \sum_{i=1}^n y_i g(x_i) \\
&= \begin{cases} f(\mathbf{x}), & \text{if } \mathbf{x} \in \{0, 1\}^n \\ +\infty, & \text{otherwise} \end{cases}
\end{aligned}$$

Therefore, $\inf_{\mathbf{x} \in [0, 1]^n} \sup_{\mathbf{y} \in \mathbb{R}^n} L(\mathbf{x}, \mathbf{y}) = \inf_{\mathbf{x} \in \{0, 1\}^n} f(\mathbf{x}) = v^*$ $\square$

A.4 Proof of Proposition 3.3

Proof. The three conditions for g derive the following properties:

- $g(\frac{1}{2}) \leq g(x) < 0$ for any $x \in (0, 1)$
- $g'(x)$ is monotone non-decreasing
- $g'(\frac{1}{2} - \delta) < 0 < g'(\frac{1}{2} + \delta)$ and $g'(\frac{1}{2} - \delta) + g'(\frac{1}{2} + \delta) = 0$

For any $t \geq 0$, the update $y_i^{t+1} - y_i^t = \beta \cdot g(x_i^t)$ lies within $[\beta \cdot g(\frac{1}{2}), 0]$. Consequently, the sequence $\{y_i^t\}_{t \geq 0}$ is monotone non-increasing. If $y_i^t > \frac{1+\Theta}{g'(\frac{1}{2}-\delta)}$ for any $t \geq 0$, there is nothing left to prove, otherwise we denote by T the smallest time index such that $y_i^T \leq \frac{1+\Theta}{g'(\frac{1}{2}-\delta)}$ and $x_i^T \notin [\frac{1}{2} - \delta, \frac{1}{2} + \delta]$. We have $y_i^T \in [2\beta \cdot g(x_i^t) + \frac{1+\Theta}{g'(\frac{1}{2}-\delta)}, \frac{1+\Theta}{g'(\frac{1}{2}-\delta)}]$, since Algorithm 1 prevents the stagnation of x_i^t within $[0.5 - \delta, 0.5 + \delta]$ for any two successive steps. Now we discuss if x_i^T lies in $[0, \frac{1}{2} - \delta]$ or $[\frac{1}{2} + \delta, 1]$.

Case I. If $x_i^T \in [0, \frac{1}{2} - \delta]$, then $g'(x_i^T) \leq g'(\frac{1}{2} - \delta) < 0$, and hence

$$\begin{aligned}\frac{\partial}{\partial x_i} L(\mathbf{x}^T, \mathbf{y}^T) &= \frac{\partial}{\partial x_i} f(\mathbf{x}^T) + y_i^T \cdot g'(x_i^T) \\ &\geq -\Theta + \frac{1 + \Theta}{g'(\frac{1}{2} - \delta)} \cdot g'(\frac{1}{2} - \delta) \\ &= 1\end{aligned}$$

For the primal update

$$\begin{aligned}x_i^{T+1} &= \Pi_{[0,1]} \left[x_i^T - \alpha \cdot \frac{\partial}{\partial x_i} L(\mathbf{x}^T, \mathbf{y}^T) \right] \\ &\leq \Pi_{[0,1]} [x_i^T - \alpha] \\ &\leq x_i^T\end{aligned}$$

Using an analogous argument, we induce that the sequence $\{x_i^{T+t}\}_{t \geq 0}$ is non-increasing and decreases at least α in each step until it reaches 0 after at most $\frac{\frac{1}{2} - \delta}{\alpha} \leq \lceil \frac{1}{2\alpha} \rceil$ steps. Therefore, we conclude that both x_i^t and y_i^t remain unchanged for $t \geq T + \lceil \frac{1}{2\alpha} \rceil$, and

$$\begin{aligned}y_i^{T + \lceil \frac{1}{2\alpha} \rceil} &= y_i^T + \sum_{t=1}^{\lceil \frac{1}{2\alpha} \rceil} (y_i^{T+t} - y_i^{T+t-1}) \\ &\geq \left(2\beta \cdot g(x_i^T) + \frac{1 + \Theta}{g'(\frac{1}{2} - \delta)} \right) + \left(\lceil \frac{1}{2\alpha} \rceil \cdot \beta \cdot g(\frac{1}{2}) \right) \\ &\geq \frac{1 + \Theta}{g'(\frac{1}{2} - \delta)} + \left(2 + \lceil \frac{1}{2\alpha} \rceil \right) \cdot \beta \cdot g(\frac{1}{2}) \quad \#g(x_i^T) \geq g(\frac{1}{2})\end{aligned}$$

Case II. If $x_i^T \in [\frac{1}{2} + \delta, 1]$, then $g'(x_i^T) \geq g'(\frac{1}{2} + \delta) > 0$, and hence

$$\begin{aligned}\frac{\partial}{\partial x_i} L(\mathbf{x}^T, \mathbf{y}^T) &= \frac{\partial}{\partial x_i} f(\mathbf{x}^T) + y_i^T \cdot g'(x_i^T) \\ &\leq \Theta + \frac{1 + \Theta}{g'(\frac{1}{2} + \delta)} \cdot g'(\frac{1}{2} + \delta) \\ &= -1\end{aligned}$$

For the primal update

$$\begin{aligned}x_i^{T+1} &= \Pi_{[0,1]} \left[x_i^T - \alpha \cdot \frac{\partial}{\partial x_i} L(\mathbf{x}^T, \mathbf{y}^T) \right] \\ &\geq \Pi_{[0,1]} [x_i^T + \alpha] \\ &\geq x_i^T\end{aligned}$$

Using an analogous argument, we induce that the sequence $\{x_i^{T+t}\}_{t \geq 0}$ is non-decreasing and increases at least α in each step until it reaches 1 after at most $\frac{\frac{1}{2} - \delta}{\alpha} \leq \lceil \frac{1}{2\alpha} \rceil$ steps. Therefore, we conclude that both x_i^t and y_i^t remain unchanged for $t \geq T + \lceil \frac{1}{2\alpha} \rceil$, and

$$\begin{aligned}y_i^{T + \lceil \frac{1}{2\alpha} \rceil} &= y_i^T + \sum_{t=1}^{\lceil \frac{1}{2\alpha} \rceil} (y_i^{T+t} - y_i^{T+t-1}) \\ &\geq \left(2\beta \cdot g(x_i^T) + \frac{1 + \Theta}{g'(\frac{1}{2} - \delta)} \right) + \left(\lceil \frac{1}{2\alpha} \rceil \cdot \beta \cdot g(\frac{1}{2}) \right) \\ &\geq \frac{1 + \Theta}{g'(\frac{1}{2} - \delta)} + \left(2 + \lceil \frac{1}{2\alpha} \rceil \right) \cdot \beta \cdot g(\frac{1}{2})\end{aligned}$$

□

A.5 Proof of Theorem 3.2

Proof. Clearly, we have the following weak duality:

$$\inf_{\mathbf{x} \in [0,1]^n} \sup_{\mathbf{y} \in \mathbb{R}^n} L(\mathbf{x}, \mathbf{y}) \geq \sup_{\mathbf{y} \in \mathbb{R}^n} \inf_{\mathbf{x} \in [0,1]^n} L(\mathbf{x}, \mathbf{y}).$$

Taking $\bar{\mathbf{y}} := \mathbf{0}$ yields $\inf_{\mathbf{x} \in [0,1]^n} L(\mathbf{x}, \bar{\mathbf{y}}) = \inf_{\mathbf{x} \in [0,1]^n} f(\mathbf{x})$. Therefore, we have

$$\begin{aligned} \sup_{\mathbf{y} \in \mathbb{R}^n} \inf_{\mathbf{x} \in [0,1]^n} L(\mathbf{x}, \mathbf{y}) &\geq \inf_{\mathbf{x} \in [0,1]^n} L(\mathbf{x}, \bar{\mathbf{y}}) \\ &= \inf_{\mathbf{x} \in [0,1]^n} f(\mathbf{x}) \\ &= \min_{\mathbf{x} \in \{0,1\}^n} f(\mathbf{x}) \quad (\text{Lemma 2.3}) \\ &= \inf_{\mathbf{x} \in [0,1]^n} \sup_{\mathbf{y} \in \mathbb{R}^n} L(\mathbf{x}, \mathbf{y}) \quad (\text{Proposition 3.1}) \end{aligned}$$

□

A.6 Proof of Theorem 3.7

Proof. Given the update $y_i^{t+1} = y_i^t + \beta \cdot g(x_i^t)$, we have $\sum_{t=0}^{T-1} -\beta \cdot g(x_i^t) = y_i^0 - y_i^T \leq y_i^0 - y_i^*$ for any $i \in [n]$ and any $T \geq 1$. Consequently,

$$\frac{\sum_{t=0}^{T-1} \sum_{i=1}^n -g(x_i^t)}{T} \leq \frac{\|\mathbf{y}^0 - \mathbf{y}^*\|_1}{\beta \cdot T}$$

Therefore, the number of Algorithm 1 to return an ϵ -binary point is bounded by $\mathcal{O}\left(\frac{\|\mathbf{y}^0 - \mathbf{y}^*\|_1}{\beta \epsilon}\right)$. □

B Generalization to the Max- k -Cut problem

Given a graph with n nodes, the Max- k -Cut problem aims to partition the nodes into k categories that maximize the number of edges connecting the nodes from distinct categories:

$$\begin{aligned} \max_{\mathbf{X} \in \mathbb{R}^{k \times n}} \quad & \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n W_{i,j} (1 - \mathbf{X}_{:,i}^\top \mathbf{X}_{:,j}) \\ \text{s.t.} \quad & \mathbf{X}_{:,i} \in \{e^{(1)}, e^{(2)}, \dots, e^{(k)}\}, \quad \forall i \in [n] \end{aligned}$$

where $\mathbf{W} \in \mathbb{R}^{n \times n}$ is a symmetric matrix representing the weights of edges and each node i is assigned with a one-hot vector $\mathbf{X}_{:,i}$ representing its category. The formulation can be equivalently expressed as:

$$\begin{aligned} \min_{\mathbf{X} \in \mathbb{R}^{k \times n}} \quad & \text{Tr}(\mathbf{X} \mathbf{W} \mathbf{X}^\top) \\ \text{s.t.} \quad & \mathbf{X}_{:,i} \in \{e^{(1)}, e^{(2)}, \dots, e^{(k)}\}, \quad \forall i \in [n] \end{aligned}$$

Note that the domain of each $\mathbf{X}_{:,i}$ can be equivalently expressed as:

$$\{e^{(1)}, e^{(2)}, \dots, e^{(k)}\} = \Delta_k \cap \{\mathbf{X}_{:,i} \in \mathbb{R}^k : g(\mathbf{X}_{:,i}) = 0\}$$

where $\Delta_k := \{\mathbf{x} \in \mathbb{R}_+^k : \sum_{i=1}^k x_i = 1\}$ denotes the k -dimensional probability simplex, and g is the convex constraint function as described in Section 3.1. We summarize two possible construction of g :

$$\bullet \quad g(\mathbf{X}_{:,i}) := \sum_{j=1}^k \mathbf{X}_{j,i}^2 - 1 \quad (\text{quadratic})$$

$$\bullet \quad g(\mathbf{X}_{:,i}) := \sum_{j=1}^k \mathbf{X}_{j,i} \log(\mathbf{X}_{j,i}) \quad (\text{entropy})$$

We can now reformulate the Max- k -Cut problem as:

$$\begin{aligned} \min_{\mathbf{X} \in \Delta_k^n} \quad & \text{Tr}(\mathbf{X} \mathbf{W} \mathbf{X}^\top) \\ \text{s.t.} \quad & g(\mathbf{X}_{:,i}) = 0 \quad \forall i \in [n] \end{aligned}$$

The Lagrangian is defined as:

$$L(\mathbf{X}, \mathbf{y}) := \text{Tr}(\mathbf{X} \mathbf{W} \mathbf{X}^\top) + \sum_{i=1}^n y_i g(\mathbf{X}_{:,i})$$

To optimize the corresponding minimax problem $\min_{\mathbf{X} \in \Delta_k^n} \max_{\mathbf{y} \in \mathbb{R}^n} L(\mathbf{X}, \mathbf{y})$, a straightforward approach is to align with Algorithm 1 by projecting $\mathbf{X}^t$ onto Δ_k^n . While there are a variety of methods for normalization, we adopt an alternative by re-parameterizing $\mathbf{X}_{:,i} \in \Delta_k$ via the softmax function:

$$\mathbf{X}_{:,i} = \text{SoftMax}(\mathbf{Z}_{:,i}) = \left(\frac{\exp(\mathbf{Z}_{j,i})}{\sum_{j=1}^k \exp(\mathbf{Z}_{j,i})} \right)_{j=1, \dots, k}$$

and the minimax problem is converted to

$$\min_{\mathbf{Z} \in \mathbb{R}^{k \times n}} \max_{\mathbf{y} \in \mathbb{R}^n} L(\text{SoftMax}(\mathbf{Z}), \mathbf{y})$$

Consequently, we can apply the approach described in Section 3.2 to the Max- k -Cut problem.

C Hyper-parameters

C.1 Experimental Settings

The multiple initial solutions are selected as follows: The initial dual variables is set to $\mathbf{y}^0 := \bar{y} \cdot \mathbf{1}$ and the initial primal variable is sampled from the uniform distribution $\mathbf{x}^0 \sim \text{Uniform}([0, 1]^n)$. The hyper-parameters can be efficiently found via a grid search on small-scale problems. Table 6 summarizes the hyper-parameters employed in our experiments.

Table 6: Hyper-parameters

Hyper-parameter	Notation	Max-Cut	MIS	Max- k -SAT	Max- k -Cut
Number of initial solutions	B	100	10	10	100
Initial dual variable	$\bar{y}$	6	5	2	6
Primal step size	α	0.025	0.02	0.01	0.01
Dual step size	β	0.025	0.02	0.005	0.01

C.2 Sensitivity Analysis

Table 7: Effect of the initial dual magnitude $\bar{y}$ on Max-Cut.

$\bar{y}$	0	1	2	4	6	8	10
G70 obj.	8756	9145 (×)	9398 (×)	9535 (✓)	9537 (✓)	9536 (✓)	9536 (✓)
G81 obj.	11338 (×)	12886 (×)	13836 (×)	13840 (✓)	13852 (✓)	13846 (✓)	13846 (✓)

Table 7 supports the landscape interpretation through an ablation on the initialization magnitude $\bar{y}$: performance improves sharply once $\bar{y}$ is large enough to satisfy the initial convexity condition.

Sensitivity to primal step size. Table 8 reports the sensitivity to α with $\beta = 0.025$ fixed. PDB0 is stable over a broad range of values, except when α is much smaller than β , in which case the landscape induced by $\mathbf{y}^t$ evolves too quickly relative to the primal update.

Table 8: Sensitivity to the primal step size α on Max-Cut, with $\beta = 0.025$.

α	0.0025	0.005	0.01	0.025	0.05	0.1	0.25
G70 obj.	9365	9451	9530	9537	9536	9535	9539
G81 obj.	13828	13842	13848	13852	13842	13854	13840

Sensitivity to dual step size. Table 9 reports the sensitivity to β with $\alpha = 0.025$ fixed. The method remains robust when the dual update is not too aggressive, while overly large β degrades solution quality. This is consistent with the iteration complexity in Theorem 3.7, which scales as $\mathcal{O}(\|\mathbf{y}^0 - \mathbf{y}^*\|_1/(\beta\epsilon))$, revealing a practical speed-quality trade-off.

Table 9: Sensitivity to the dual step size β on Max-Cut, with $\alpha = 0.025$.

β	0.001	0.0025	0.005	0.01	0.025	0.05	0.1	0.25
G70 obj.	9538	9536	9542	9534	9537	9509	9495	9415
G81 obj.	13880	13868	13862	13858	13852	13840	13724	13480

Sensitivity to batch size. Table 10 reports the sensitivity to the number of parallel initializations B . Larger batches improve objectives at the cost of runtime, while the overall sensitivity remains modest.

Table 10: Sensitivity to batch size B on Max-Cut. Runtime is reported in seconds.

B	1	10	100	1000
G70 obj. (time)	9525 (0.6)	9527 (0.7)	9537 (1.9)	9540 (9.6)
G81 obj. (time)	13828 (0.9)	13842 (1.1)	13852 (3.3)	13858 (35.8)

D Additional Experimental Results

D.1 Complete results on Gset instances for Max-Cut and Max-k-cut

We provide complete results on the Gset instances. Table 11 exhibits the complete results on Gset instances for the Max-Cut problem, while Table 12 exhibits the complete results on Gset instances for the Max-3-Cut problem.

Notably, as PDBO exhibits rapid convergence, its solutions can be effectively refined by using them as initial points for ABS2—a method that combines multiple local search strategies and demonstrates strong performance given sufficient computation time. This hybrid approach, denoted as PDBO+ABS2, leverages the fast convergence of PDBO with the refinement capability of ABS2.

While PDBO consistently delivers superior performance on Max-3-Cut, the highly specialized ABS2 solver is a strong performer on small-scale Max-Cut instances. On the 66 benchmarks with up to 10,000 nodes, PDBO alone is competitive, matching or beating ABS2 in 16 cases. The most effective strategy, however, is a hybrid approach: using PDBO for rapid initialization followed by ABS2’s intensive local search. This hybrid method matches or surpasses standalone ABS2 on 44 of the 66 instances, demonstrating a valuable synergy.

Moreover, note that ABS2 is not applicable to Max-k-Cut and is confined to QUBO/Max-Cut. In contrast, PDBO’s first-order optimization strategy allows it to address a broader class of combinatorial problems with significant efficiency.

Table 11: Computational results on Gset instances for Max-Cut.

Instance	V	E	PIGMW		CRA		ROS		ANYCSP		FEM		ABS2		PDB0		PDB0+ABS2	
			Obj. ↑	Time (s) ↓	Obj. ↑	Time (s) ↓	Obj. ↑	Time (s) ↓	Obj. ↑	Time (s) ↓	Obj. ↑	Time (s) ↓	Obj. ↑	Time (s) ↓	Obj. ↑	Time (s) ↓	Obj. ↑	Time (s) ↓
G1	800	19176	10891	14.3	11398	25.0	11313	0.8	11623	6.5	11624	0.8	11624	0.1	11624	1.6	11624	1.6
G2	800	19176	10930	23.3	11498	23.7	11417	1.1	11612	10.5	11617	0.8	11620	0.2	11617	1.4	11620	3.6
G3	800	19176	10958	0.2	11464	21.5	11383	0.7	11620	3.8	11622	0.8	11622	0.2	11621	1.0	11622	1.2
G4	800	19176	10973	0.7	11486	24.0	11357	0.6	11644	9.8	11646	0.7	11646	1.1	11646	0.9	11646	0.9
G5	800	19176	10901	0.2	11463	24.8	11410	0.8	11630	7.3	11627	0.7	11631	0.1	11631	1.7	11631	1.7
G6	800	19176	1751	33.3	1747	46.7	1949	0.9	2173	13.5	2178	0.8	2178	0.1	2178	1.4	2178	1.4
G7	800	19176	1548	19.0	1512	48.1	1772	0.9	1997	3.7	2001	0.8	2006	7.1	2001	2.2	2001	2.2
G8	800	19176	1588	58.5	1490	48.1	1713	1.1	2004	10.2	2004	0.8	2005	0.2	2003	1.3	2005	5.6
G9	800	19176	1621	26.8	1544	46.8	1784	0.8	2044	13.5	2051	0.8	2054	0.1	2054	1.4	2054	1.4
G10	800	19176	1635	21.0	1481	48.1	1755	1.3	1991	4.8	1999	0.8	2000	3.0	1999	0.9	1999	0.9
G11	800	1600	478	3.1	440	45.2	492	0.7	558	1.4	558	0.3	564	0.1	562	0.6	564	1.7
G12	800	1600	472	3.6	428	44.7	476	0.4	548	0.3	552	0.4	556	0.5	554	0.5	554	0.5
G13	800	1600	490	4.6	430	44.2	514	0.6	572	0.4	576	0.3	582	0.5	582	0.5	582	0.5
G14	800	4694	2815	12.3	2987	39.2	2941	0.7	3058	6.2	3058	0.5	3064	2.1	3054	0.8	3064	22.9
G15	800	4661	2746	3.8	2951	38.4	2944	0.7	3040	2.8	3047	0.5	3045	1.2	3041	0.9	3043	2.6
G16	800	4672	2735	1.8	2966	39.5	2955	0.8	3040	3.4	3049	0.5	3051	1.2	3044	0.9	3046	79.9
G17	800	4667	2699	40.8	2948	39.2	2952	0.8	3037	2.2	3043	0.5	3042	63.6	3037	0.8	3041	8.5
G18	800	4694	833	14.9	822	44.9	859	1.1	989	4.2	991	0.5	991	33.0	987	0.8	992	8.1
G19	800	4661	721	19.9	771	45.8	776	0.8	899	1.5	905	0.5	906	0.7	902	0.7	902	0.7
G20	800	4672	668	7.7	771	44.3	838	1.2	936	1.7	941	0.5	941	0.1	940	1.0	941	1.2
G21	800	4667	753	11.1	766	46.7	810	0.5	925	5.7	930	0.4	931	1.1	924	0.7	931	2.5
G22	2000	19990	12308	13.1	13116	35.4	12920	1.0	13347	13.9	13356	0.8	13357	10.4	13347	2.4	13359	4.8
G23	2000	19990	12288	2.1	13109	35.0	12996	1.4	13328	7.5	13338	0.8	13334	1.4	13327	2.3	13337	41.6
G24	2000	19990	12283	12.1	13067	35.8	12935	1.1	13316	9.9	13329	0.8	13327	32.9	13317	2.4	13317	2.4
G25	2000	19990	12251	12.6	13122	36.2	13008	0.6	13323	12.9	13334	0.8	13329	55.2	13321	2.2	13325	18.3
G26	2000	19990	12247	14.4	13028	36.4	12952	0.8	13311	10.8	13319	0.8	13314	4.7	13309	2.3	13315	4.1
G27	2000	19990	2576	57.2	2749	46.6	2895	1.0	3314	16.2	3339	0.8	3331	8.0	3326	1.1	3326	1.1
G28	2000	19990	2581	122.1	2685	46.6	2972	0.8	3275	18.1	3291	0.8	3283	1.5	3296	2.1	3296	2.1
G29	2000	19990	2730	50.6	2763	46.6	2991	1.1	3381	17.6	3396	0.8	3391	36.4	3383	2.4	3404	23.9
G30	2000	19990	2675	130.2	2748	46.8	2951	1.0	3399	6.2	3409	0.8	3411	0.9	3402	1.9	3402	1.9
G31	2000	19990	2601	40.4	2641	47.6	2890	0.7	3296	16.3	3302	0.8	3288	1.0	3302	2.3	3306	5.6
G32	2000	4000	1170	5.9	1156	44.7	1234	1.3	1386	8.7	1388	0.5	1410	72.6	1398	0.6	1410	61.7
G33	2000	4000	1134	12.0	1122	43.8	1220	0.7	1354	3.3	1360	0.4	1382	52.7	1370	0.6	1380	94.5
G34	2000	4000	1122	14.4	1124	44.8	1204	0.5	1358	7.0	1358	0.4	1384	9.9	1376	0.9	1384	127.6
G35	2000	11778	6968	4.5	7453	39.4	7460	0.7	7650	8.0	7666	0.7	7661	12.1	7643	1.2	7660	8.2

Instance	$ \gamma $	$ \mathcal{E} $	PIGNN		CRA		ROS		ANYCSP		FEM		ABS2		PDBO		PDBO+ABS2	
			Obj. $\uparrow$	Time (s) $\downarrow$	Obj. $\uparrow$	Time (s) $\downarrow$	Obj. $\uparrow$	Time (s) $\downarrow$	Obj. $\uparrow$	Time (s) $\downarrow$	Obj. $\uparrow$	Time (s) $\downarrow$	Obj. $\uparrow$	Time (s) $\downarrow$	Obj. $\uparrow$	Time (s) $\downarrow$	Obj. $\uparrow$	Time (s) $\downarrow$
G36	2000	11766	6967	5.7	7435	38.9	7423	1.4	7644	11.7	7659	0.7	7666	31.4	7636	1.0	7646	5.6
G37	2000	11785	7006	6.7	7466	39.9	7432	1.2	7654	8.9	7670	0.7	7667	56.8	7676	0.9	7676	54.1
G38	2000	11779	6993	1.3	7450	39.1	7468	1.8	7652	7.6	7664	0.7	7672	14.2	7654	1.5	7665	13.6
G39	2000	11778	1961	12.3	2115	46.8	2113	1.2	2384	6.8	2396	0.6	2385	1.4	2387	2.0	2387	2.0
G40	2000	11766	1915	104.4	2112	44.5	2149	1.6	2376	13.9	2394	0.7	2389	3.7	2365	1.4	2396	83.8
G41	2000	11785	1976	21.1	2108	46.0	2090	1.8	2371	5.3	2391	0.6	2405	33.4	2369	1.5	2376	3.3
G42	2000	11779	1972	31.3	2203	44.4	2191	0.6	2462	3.8	2464	0.7	2468	2.4	2450	1.3	2476	87.8
G43	1000	9990	6210	22.9	6537	33.8	6427	0.8	6660	7.3	6660	0.6	6660	0.1	6660	1.2	6660	1.2
G44	1000	9990	6160	1.0	6546	33.9	6433	0.6	6646	3.3	6650	0.6	6650	3.3	6650	1.3	6650	1.3
G45	1000	9990	6152	1.7	6520	35.4	6473	0.8	6646	2.0	6653	0.6	6654	7.3	6653	1.1	6653	1.1
G46	1000	9990	6122	2.3	6564	34.1	6463	0.8	6639	7.7	6646	0.6	6649	2.0	6649	1.0	6649	1.0
G47	1000	9990	6174	1.6	6503	35.0	6487	0.8	6655	8.9	6656	0.6	6657	3.2	6656	1.6	6657	11.2
G48	3000	6000	5072	23.1	5852	41.9	5624	0.9	6000	1.5	6000	0.3	6000	0.2	6000	0.4	6000	0.4
G49	3000	6000	4992	16.0	5898	42.1	5552	0.9	6000	0.9	6000	0.3	6000	0.3	6000	0.4	6000	0.4
G50	3000	6000	5040	16.7	5756	42.7	5568	0.8	5878	2.1	5880	0.4	5880	0.4	5880	1.1	5880	1.1
G51	1000	5909	3407	2.3	3732	39.3	3711	1.0	3832	5.8	3843	0.5	3845	31.3	3834	0.9	3841	5.5
G52	1000	5916	3458	1.1	3742	39.2	3734	1.7	3837	7.0	3845	0.6	3846	70.9	3834	0.8	3841	4.8
G53	1000	5914	3526	9.3	3739	39.3	3736	1.5	3836	2.5	3842	0.6	3847	31.5	3838	0.9	3840	2.2
G54	1000	5916	3314	2.4	3743	39.4	3722	1.4	3838	7.2	3844	0.5	3845	35.2	3836	1.1	3851	174.6
G55	5000	12498	9059	84.3	10008	43.1	9797	1.4	10248	22.6	6418	0.2	10254	93.5	10240	1.3	10258	120.3
G56	5000	12498	3127	106.0	3414	47.1	3439	0.9	3978	13.7	127	0.2	3970	63.0	3956	1.5	3970	136.7
G57	5000	10000	2794	18.5	2942	45.9	3036	0.9	3398	18.0	3408	0.8	3474	150.0	3460	1.2	3478	157.7
G58	5000	29570	17435	7.8	18695	41.8	18692	1.4	19187	29.7	19211	1.3	19200	120.2	19163	3.6	19192	180.2
G59	5000	29570	4734	170.8	5580	48.6	5298	0.9	5978	33.1	6036	1.4	6007	51.8	5967	2.7	6006	98.2
G60	7000	17148	12382	32.4	13764	46.1	13413	1.0	14108	27.3	8746	0.2	14138	171.1	14111	2.0	14111	2.0
G61	7000	17148	4518	58.0	4951	49.4	5061	1.0	5736	19.9	326	0.2	5707	30.7	5717	1.8	5717	1.8
G62	7000	14000	3878	16.0	4108	48.6	4268	1.3	4758	13.6	4762	1.1	4848	176.5	4808	1.8	4828	175.2
G63	7000	41459	24237	106.0	26183	43.5	26209	2.3	26844	48.0	26941	2.0	26893	179.4	26881	3.9	26881	3.9
G64	7000	41459	6829	105.0	8114	50.3	7665	1.3	8599	45.9	8660	1.9	8622	14.9	8591	4.5	8591	4.5
G65	8000	16000	4446	35.2	4814	50.9	4898	1.0	5414	26.9	5420	1.6	5520	148.3	5482	1.6	5504	167.9
G66	9000	18000	5142	20.9	5488	51.7	5598	1.5	6192	38.2	6196	2.0	6302	165.9	6268	1.8	6318	160.2
G67	10000	20000	5538	23.7	5948	53.7	6144	1.5	6772	39.9	6782	2.4	6880	156.3	6872	2.5	6894	175.8
G70	10000	9999	8534	25.2	9240	53.9	8872	1.8	9379	35.7	9510	0.2	9510	175.5	9537	1.9	9537	1.9
G72	10000	20000	5588	44.5	6058	53.9	6148	1.2	6816	36.1	6824	2.6	6932	172.2	6906	2.0	6950	170.9
G77	14000	28000	7896	42.1	8720	75.9	8746	2.2	9686	53.5	9688	4.0	9824	171.1	9812	2.1	9840	5.1
G81	20000	40000	11078	157.6	12450	120.4	12320	5.2	13670	73.5	13684	7.5	13850	177.1	13852	3.3	13860	4.6

Table 12: Computational results on Gset instances for Max-3-Cut.

Instance	$ \mathcal{V} $	$ \mathcal{E} $	ROS		ANYCSP		FEM		PDBO	
			Obj. $\uparrow$	Time (s) $\downarrow$	Obj. $\uparrow$	Time (s) $\downarrow$	Obj. $\uparrow$	Time (s) $\downarrow$	Obj. $\uparrow$	Time (s) $\downarrow$
G1	800	19176	14892	1.2	15121	18.0	15058	1.0	15144	3.1
G2	800	19176	14892	1.3	15119	15.8	15070	0.8	15159	3.7
G3	800	19176	14932	1.2	15129	17.0	15080	1.1	15161	4.2
G4	800	19176	14883	1.3	15141	15.7	15080	0.9	15182	4.5
G5	800	19176	14859	1.4	15144	16.8	15087	1.1	15185	4.6
G6	800	19176	2341	1.2	1364	10.5	2509	0.8	2616	2.9
G7	800	19176	2152	1.5	1131	0.2	2306	0.8	2393	3.9
G8	800	19176	2083	1.6	1114	0.2	2321	1.1	2413	3.6
G9	800	19176	2214	1.5	1264	0.2	2357	0.8	2449	3.2
G10	800	19176	2104	1.4	1110	0.2	2297	0.9	2393	4.1
G11	800	1600	607	0.8	655	2.8	645	0.3	660	1.9
G12	800	1600	600	0.8	646	5.6	633	0.3	653	1.9
G13	800	1600	621	0.8	671	5.4	661	0.3	677	1.7
G14	800	4694	3886	0.9	3987	7.3	3960	0.4	3998	2.2
G15	800	4661	3855	1.0	3966	6.5	3933	1.1	3968	2.1
G16	800	4672	3876	0.9	3966	7.3	3934	0.4	3971	2.2
G17	800	4667	3861	0.8	3956	3.9	3929	1.1	3964	2.0
G18	800	4694	1080	1.0	1110	7.0	1140	0.3	1188	2.2
G19	800	4661	939	0.9	971	7.8	1020	0.6	1062	2.3
G20	800	4672	966	0.9	1004	7.7	1048	0.3	1106	2.2
G21	800	4667	978	0.9	1010	5.7	1063	0.4	1102	2.0
G22	2000	19990	16679	1.5	17068	16.0	16964	0.8	17122	4.3
G23	2000	19990	16676	1.3	17084	25.2	16954	0.8	17127	4.3
G24	2000	19990	16701	1.4	17063	22.1	16959	1.0	17119	4.8
G25	2000	19990	16690	1.4	17072	24.2	16971	1.0	17115	4.2
G26	2000	19990	16686	1.4	17056	22.2	16941	0.9	17113	4.1
G27	2000	19990	3534	1.4	3129	24.9	3814	0.8	3968	4.2
G28	2000	19990	3456	1.4	3030	13.2	3752	0.8	3921	4.1
G29	2000	19990	3593	2.1	3225	9.3	3899	2.4	4056	4.9
G30	2000	19990	3593	1.6	3231	14.9	3887	2.1	4078	5.6
G31	2000	19990	3434	1.4	3067	18.9	3773	0.8	3956	5.8
G32	2000	4000	1509	1.1	1602	11.3	1591	0.3	1635	2.3
G33	2000	4000	1446	1.0	1572	11.8	1553	1.1	1602	2.3
G34	2000	4000	1437	1.0	1551	11.1	1546	0.4	1591	2.3
G35	2000	11778	9752	1.2	9981	16.3	9897	1.9	9997	3.4
G36	2000	11766	9744	1.1	9978	18.2	9873	1.9	10003	3.5
G37	2000	11785	9766	1.1	9989	17.0	9902	1.8	10004	3.1
G38	2000	11779	9795	1.2	9986	4.8	9893	1.9	10000	3.1
G39	2000	11778	2533	1.4	2525	9.7	2722	0.5	2857	3.8
G40	2000	11766	2461	1.3	2536	16.3	2671	1.9	2833	3.3
G41	2000	11785	2524	1.2	2546	17.8	2679	0.5	2834	3.8
G42	2000	11779	2582	1.2	2631	18.0	2777	1.9	2926	3.4
G43	1000	9990	8312	1.1	8525	4.3	8478	0.6	8566	3.3
G44	1000	9990	8349	1.4	8523	12.8	8486	0.6	8550	3.1
G45	1000	9990	8354	1.6	8512	11.9	8486	0.5	8559	3.0
G46	1000	9990	8346	4.1	8508	7.2	8473	1.5	8558	3.0
G47	1000	9990	8338	2.7	8523	10.7	8483	0.5	8569	3.1
G48	3000	6000	5920	2.7	6000	7.7	6000	0.4	6000	2.3
G49	3000	6000	5914	2.3	6000	6.9	6000	0.4	6000	2.4
G50	3000	6000	5924	1.9	5998	9.7	5998	0.5	6000	2.7
G51	1000	5909	4897	1.0	5004	8.3	4957	0.4	5012	2.3
G52	1000	5916	4910	1.0	5019	9.1	4970	0.4	5012	2.3
G53	1000	5914	4918	1.0	5007	4.7	4964	0.4	5015	2.2
G54	1000	5916	4886	0.9	5015	8.7	4966	0.4	5013	2.2
G55	5000	12498	11992	1.4	12351	27.3	12210	2.2	12345	4.3
G56	5000	12498	4126	1.7	4495	31.3	4425	2.2	4663	4.0
G57	5000	10000	3688	1.5	3940	29.4	3900	0.4	4028	3.6
G58	5000	29570	24482	4.2	25027	46.5	24807	3.7	25088	5.4
G59	5000	29570	6549	5.2	6316	43.3	6785	3.9	7190	8.0
G60	7000	17148	16496	1.9	16978	41.6	16797	2.9	16961	4.5
G61	7000	17148	5989	2.0	6464	41.2	6401	2.9	6730	4.8
G62	7000	14000	5149	1.9	5487	37.0	5441	2.5	5635	4.3
G63	7000	41459	34331	3.0	35087	49.3	34776	5.3	35194	8.0
G64	7000	41459	9399	2.8	9219	48.1	9767	5.1	10333	8.4
G65	8000	16000	5914	2.2	6255	44.3	6207	2.9	6425	4.5
G66	9000	18000	6739	2.5	7147	39.4	7103	3.2	7338	4.5
G67	10000	20000	7364	3.0	7797	55.9	7748	3.3	8015	6.0
G70	10000	9999	9983	1.9	9909	16.2	9999	1.4	9999	3.3
G72	10000	20000	7435	2.9	7906	58.9	7835	0.7	8111	4.7
G77	14000	28000	10559	5.3	11158	84.0	11102	4.4	11467	5.0
G81	20000	40000	14907	9.7	15727	115.0	15683	6.1	16191	7.9

D.2 Experiments for GPU efficiency

In this section we conduct comparison between the CPU/GPU versions of PDBO.




Figure 3: The objective value of Max-Cut on G67, as a function of runtime

Figure 3 compares the solution quality of each method throughout the optimization process, plotting the objective value against runtime under a 180-second time limit, on the G67 instance for Max-Cut. Clearly, PDBO-GPU (blue) performs much faster than PDBO-CPU (red). The initial latency of PDBO-GPU could be attributed to the *just-in-time* nature of Jax.

D.3 Experiments on MIS

Given a graph $G := (V, E)$, the MIS problem seeks the largest vertex subset $S \subseteq V$ such that no two vertices in S are adjacent. This can be formulated as:

$$\begin{aligned} \max_{\mathbf{x} \in \{0,1\}^n} \quad & \sum_{i=1}^n x_i \\ \text{s.t.} \quad & x_i x_j = 0, \quad \forall (i, j) \in E \end{aligned}$$

where $x_i \in \{0, 1\}$ indicates whether vertex i is included in the independent set. Following the approach of [17, 18], we reformulate MIS as a QUBO by penalizing the adjacency constraints:

$$\max_{\mathbf{x} \in \{0,1\}^n} \sum_{i=1}^n x_i - \lambda \cdot \sum_{(i,j) \in E} x_i x_j$$

where $\lambda > 0$ is a penalty parameter. In line with established practice, we set $\lambda = 4$ for all methods to ensure a consistent comparison.

For MIS, we employ d -regular random graphs (d -RRGs), where each node has degree d , following the experimental setup of [17]. To evaluate performance across both sparse and dense graph structures, we consider degrees $d \in \{3, 100\}$ and graph sizes $n \in \{10^4, 5 \times 10^4\}$ nodes. For each (n, d) configuration, we generate 20 independent graph instances to ensure statistical significance.

The objective of the MIS problem is to find the largest set of non-adjacent nodes. Although we formulate MIS as a QUBO via constraint penalization, we only report the objective value (independent set size) for solutions that satisfy the independence constraints. We evaluate PDBO against baseline algorithms PIGNN, CRA, FEM, ABS2 and Gurobi on d -RRG instances, with performance averaged over 20 instances per configuration, as shown in Table 13. The results show that PIGNN, CRA, FEM and even Gurobi struggle on dense graphs with degree 100 (entries marked with “–” indicate that FEM failed to return feasible solutions). This aligns with the known computational hardness of dense graphs [30]. In contrast, both PDBO and ABS2 achieve strong performance even on dense instances. Notably, PDBO outperforms all baselines, identifying better solutions in less time on both sparse and dense graphs.

Table 13: Results on d -RRGs for MIS

Method	$(10^4, 3)$		$(10^4, 100)$		$(5 \times 10^4, 3)$		$(5 \times 10^4, 100)$	
	Obj $\uparrow$	Time $\downarrow$	Obj $\uparrow$	Time $\downarrow$	Obj $\uparrow$	Time $\downarrow$	Obj $\uparrow$	Time $\downarrow$
PIGNN	4176.1 ± 19.3	66.7 ± 14.2	7.5 ± 6.7	0.2 ± 0.1	20705.7 ± 47.1	178.9 ± 1.0	15.8 ± 13.7	0.6 ± 0.5
CRA	4361.4 ± 8.1	52.2 ± 0.1	361.9 ± 187.6	46.6 ± 23.2	166.0 ± 75.7	20.4 ± 4.0	11.1 ± 9.8	0.5 ± 0.5
FEM	4417.2 ± 4.4	0.7 ± 0.0	—	180.0 ± 0.0	22053.8 ± 11.1	1.3 ± 0.1	—	180.0 ± 0.0
ABS2	4394.8 ± 6.5	154.2 ± 27.6	587.4 ± 3.5	140.0 ± 36.2	21307.7 ± 43.6	173.1 ± 5.7	2971.5 ± 12.8	175.0 ± 3.9
Gurobi	4141.8 ± 19.4	147.2 ± 24.8	448.0 ± 5.1	0.1 ± 0.0	20144.3 ± 181.7	96.6 ± 49.9	2235.6 ± 12.7	0.1 ± 0.0
PDBO	4431.9 ± 3.1	0.6 ± 0.0	603.0 ± 3.3	3.6 ± 0.3	22127.3 ± 10.3	1.1 ± 0.1	3006.2 ± 6.2	14.6 ± 0.6

D.4 Experiments on MAP

In this section we conduct additional experiments on the Maximum a Posterior (MAP) inference in discrete Markov random fields (MRF). Let $\mathcal{G}$ be a graph of n nodes and $\mathcal{C}$ is the set of cliques, and let $\mathcal{S} = \mathcal{S}_1 \times \dots \times \mathcal{S}_n$ be the assignment domain of the n nodes (i.e. variables). Consider an MRF representing a joint distribution

$$p(s) = \frac{1}{Z} \prod_{C \in \mathcal{C}} \psi_C(s_C), \quad \forall s \in \mathcal{S}$$

where s_C is the joint assignment of the variables in clique C , ψ_C are positive functions called potentials, and Z is the partition function.

The MAP inference problem is to find the most likely assignment to the variables, which can be equivalently modeled as:

$$\min_{s \in \mathcal{S}} \sum_{C \in \mathcal{C}} -\log \psi_C(s_C)$$

As shown by [31], when $|\mathcal{S}_i| = 2$ for all i , the MAP inference can be reformulated as an binary optimization with polynomial objective, which is applicable to PDBO. Furthermore, one can easily handles the cases where $|\mathcal{S}_i| > 2$ with similar techniques as we discussed in Appendix B.

D.4.1 Protein Prediction

In particular, we adopt the protein-prediction dataset [32], which is widely adopted in the related literature and competitions [33, 34]. The dataset includes 8 instances and Table 14 presents the average objective value and running time for each method. Please refer to [33] for detailed description on the baselines as well as the formal problem definition.

Table 14: Results for MAP on the protein-prediction dataset

	Obj. $\downarrow$	Time $\downarrow$
ogm-LBP-LF2	52942.95	69.86
ogm-LF-3	57944.06	25.99
ogm-LBP-0.5	53798.89	60.97
ogm-TRBP-0.5	61386.17	86.03
ADDD	106216.86	10.97
MPLP	101531.75	86.6
ogm-LP-LP	102918.41	180.72
ogm-ILP	60047.02	2262.83
BRAOBB-1	61079.07	3600.16
PDBO	52842.61	15.07

The results show that PDBO is both efficient and effective for handling the MAP problems.

D.4.2 UAI2022 Competition

We also evaluate PDBO on the datasets from the UAI 2022 competition (<https://uaicompetition.github.io/uci-2022/>), a standard modern benchmark. In particular, we consider the **Grids** dataset

(7 instances in total) that aligns with the binary optimization formulation (1). We compare against three state-of-the-art baselines—LBP, TRBP, and Toulbar2—with their results taken directly from the recent and comprehensive study of [35]:

- Loopy Belief Propagation (denoted as LBP) & Tree-reweighted Belief Propagation (denoted as TRBP) [36]: Classic and widely-used message-passing algorithms.
- Toulbar2 [37]: the winner on all MPE and MMAP task categories of UAI 2022 Inference Competition.

Table 15: Results on the UAI 2022 benchmarks, lower values imply better performance

	Grids_19	Grids_21	Grids_24	Grids_25	Grids_26	Grids_27	Grids_30
LBP	-2250.440	-13119.300	-13210.400	-2170.890	-2063.350	-9024.640	-2142.890
TRBP	-2103.610	-12523.300	-13260.900	-2171.050	-1903.910	-9019.470	-2154.910
Toulbar2	-2643.107	-18895.393	-18274.302	-2620.268	-3010.719	-12284.284	-2984.248
PDBO	-2696.341	-19235.873	-18691.510	-2653.470	-3021.883	-12481.272	-2989.997

The results are shown in Table 15. PDBO achieves the best performance on all seven instances. Notably, it finds these superior solutions within 5.9 seconds on average, while the baselines from [35] (including Toulbar2 with a 1,200-second limit) could not match this performance. These results, conducted on a standard modern benchmark and compared against baseline results, robustly demonstrate PDBO’s effectiveness and efficiency.

D.5 Dual Certificate

When the initial value $\mathbf{y}^0$ is sufficiently large (e.g. $\bar{y} \geq -\lambda_{\min}(W)$ for max-cut), the Lagrangian extension $f_{\mathbf{y}}(\mathbf{x}) := L(\mathbf{x}, \mathbf{y})$ turns out to be convex in $\mathbf{x}$. In this regime, we are able to globally minimize $f_{\mathbf{y}}(\mathbf{x})$ via classic gradient descent (albeit PDBO employs a single gradient descent step instead), the corresponding optimal value would provide a valid dual bound on the optimal objective value, serving as a dual certificate. We present the corresponding dual bounds of the Max-Cut problem implied by PDBO in Table 16.

Table 16: Dual bounds for Max-Cut, implied by PDBO

Max-Cut	G67	G70	G72	G77	G81
primal bound	6872	9537	6906	9812	13852
dual bound	8843.7	13891.5	8906.1	12767.7	18172.2